

Kubernetes lab

Exercise: Pods

Pods are the smallest, most basic deployable objects in Kubernetes. A Pod represents a single instance of a running process in your cluster. Pods contain one or more containers, such as Docker containers. Although you want deploy pods directly (static pods), knowledge for defining pods

manifest files will be used for defining more complex Kubernetes resources like Controllers.

Practice1: Simple pods operations

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so.

Copy/paste will not help you to learn Kubernetes!

1. Login to Azure and connect to your AKS cluster.

The screenshot displays the Microsoft Azure portal interface. The top navigation bar shows the user is logged in as 'Stanislav Nikolov New1@...'. The main content area is titled 'clusterHomework' and shows the 'Overview' tab. The 'Essentials' section provides a quick overview of the cluster's status and configuration. The 'Properties' tab is selected, showing detailed information about the cluster's configuration, including networking, node pools, and extensions.

| Section | Property | Value |
|---------------|----------------------------------|---------------------------------------------------|
| Essentials | Resource group | clusterHomework_group |
| | Status | Succeeded (Running) |
| | Location | East US |
| | Subscription | Azure Pass - Sponsorship |
| Networking | API server address | clusterhomework-dns-222g5ort.hcp.eastus.azmk8s.io |
| | Network type (plugin) | Kubenet |
| | Pod CIDR | 10.244.0.0/16 |
| | Service CIDR | 10.0.0.0/16 |
| Node pools | Node pools | 1 node pool |
| | Kubernetes versions | 1.24.10 |
| | Node sizes | Standard_DS2_v2 |
| | Encryption type | Encryption at-rest with a platform-managed key |
| Configuration | Kubernetes version | 1.24.10 |
| | Auto Upgrade Type | Patch |
| | Authentication and Authorization | Local accounts with Kubernetes RBAC |
| | Local accounts | Enabled |

2. Check how many pods run under the default namespace.
Run `kubectl get pods`.

```
VERBOSE: Building your Azure drive ...  
PS /home/stanislav> az account set --subscription c983dec5-cde0-4991-9469-c26f8cf60056  
PS /home/stanislav> az aks get-credentials --resource-group clusterHomework_group --name clusterHomework  
Merged "clusterHomework" as current context in /home/stanislav/.kube/config  
PS /home/stanislav> kubectl run nginx --image=nginx  
pod/nginx created  
PS /home/stanislav>
```

3. You should not see any pod under the default namespace. Now check all namespaces. Run `kubectl get pods --all-namespaces`.

```
PS /home/stanislaw> kubectl get pods --all-namespaces
```

| NAMESPACE | NAME | READY | STATUS | RESTARTS | AGE |
|-------------|-------------------------------------|-------|---------|----------|-------|
| kube-system | ama-logs-7bnjj | 2/2 | Running | 0 | 9m47s |
| kube-system | ama-logs-c75mm | 2/2 | Running | 0 | 9m51s |
| kube-system | ama-logs-r77cl | 2/2 | Running | 0 | 9m52s |
| kube-system | ama-logs-rs-79b79975d5-5dgmr | 1/1 | Running | 0 | 9m58s |
| kube-system | azure-ip-masq-agent-6nbj8 | 1/1 | Running | 0 | 9m51s |
| kube-system | azure-ip-masq-agent-m7pjm | 1/1 | Running | 0 | 9m52s |
| kube-system | azure-ip-masq-agent-nlzh4 | 1/1 | Running | 0 | 9m47s |
| kube-system | cloud-node-manager-5qbwg | 1/1 | Running | 0 | 9m51s |
| kube-system | cloud-node-manager-rtrnm | 1/1 | Running | 0 | 9m47s |
| kube-system | cloud-node-manager-zbh2z | 1/1 | Running | 0 | 9m52s |
| kube-system | coredns-59b6bf8b4f-dcxbv | 1/1 | Running | 0 | 9m58s |
| kube-system | coredns-59b6bf8b4f-vrnp2 | 1/1 | Running | 0 | 8m37s |
| kube-system | coredns-autoscaler-64b6477b8b-njsnl | 1/1 | Running | 0 | 9m58s |
| kube-system | csi-azuredisk-node-dnf6r | 3/3 | Running | 0 | 9m52s |
| kube-system | csi-azuredisk-node-qz4ks | 3/3 | Running | 0 | 9m47s |
| kube-system | csi-azuredisk-node-xl5rx | 3/3 | Running | 0 | 9m51s |
| kube-system | csi-azurefile-node-5dmrq | 3/3 | Running | 0 | 9m51s |
| kube-system | csi-azurefile-node-br9w4 | 3/3 | Running | 0 | 9m52s |
| kube-system | csi-azurefile-node-nwgdz | 3/3 | Running | 0 | 9m47s |
| kube-system | konektivitiy-agent-7cf8bd6556-b7828 | 1/1 | Running | 0 | 9m58s |
| kube-system | konektivitiy-agent-7cf8bd6556-t6qqh | 1/1 | Running | 0 | 9m58s |
| kube-system | kube-proxy-n8q64 | 1/1 | Running | 0 | 9m47s |
| kube-system | kube-proxy-tg9j6 | 1/1 | Running | 0 | 9m51s |
| kube-system | kube-proxy-v52qk | 1/1 | Running | 0 | 9m52s |
| kube-system | metrics-server-5f8d84558d-d7cbb | 2/2 | Running | 0 | 8m30s |
| kube-system | metrics-server-5f8d84558d-xnpnd | 2/2 | Running | 0 | 8m30s |

```
PS /home/stanislaw>
```

4. How many pods do you see? Who deployed these pods? Why are they deployed?

- There are kube-system pods mostly. Around 28

5. Now deploy you first pod using the imperative approach. Run `kubectl run nginx --image=nginx`.

```
PS /home/stanislaw> kubectl run nginx --image=nginx
pod/nginx created
PS /home/stanislaw>
```

6. Validate if the pods has been created. What is the status of your pod?

```
PS /home/stanislaw> kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           83s
PS /home/stanislaw>
```

7. Check the logs coming out of your pod. Run kubectl logs nginx.

```
PS /home/stanislaw> kubectl logs nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2023/04/09 07:38:08 [notice] 1#1: using the "epoll" event method
2023/04/09 07:38:08 [notice] 1#1: nginx/1.23.4
2023/04/09 07:38:08 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2023/04/09 07:38:08 [notice] 1#1: OS: Linux 5.4.0-1104-azure
2023/04/09 07:38:08 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2023/04/09 07:38:08 [notice] 1#1: start worker processes
2023/04/09 07:38:08 [notice] 1#1: start worker process 29
2023/04/09 07:38:08 [notice] 1#1: start worker process 30
PS /home/stanislaw>
```

8. Run following command to check current resource consumption of your pod: kubectl top pod nginx.

```
PS /home/stanislaw> kubectl top pod nginx
NAME      CPU(cores)   MEMORY(bytes)
nginx     0m           3Mi
PS /home/stanislaw>
```

9. Check on which Node your pods has been scheduled. Run kubectl get pods -o wide.

```
PS /home/stanislaw> kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE                                     NOMINATED NODE   READINESS GATES
nginx     1/1     Running   0           5m9s  10.244.0.8   aks-agentpool-35627940-vmss000000    <none>           <none>
PS /home/stanislaw>
```

10. Try to find the same information but this time running kubectl describe pod nginx.

```
Service Account: default
Node: aks-agentpool-35627940-vmss000000/10.244.0.6
Start Time: Sun, 09 Apr 2023 07:38:04 +0000
Labels: run=nginx
Annotations: <none>
Status: Running
IP: 10.244.0.8
IPs:
  IP: 10.244.0.8
Containers:
  nginx:
    Container ID: containerd://12dea4222105c5e3005ba5d167c0f58555ef343d604ae7ec7e10561226741831
    Image: nginx
    Image ID: docker.io/library/nginx@sha256:2ab30d6ac53580a6db8b657abf0f68d75360ff5cc1670a85acb5bd85ba1b19c0
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Sun, 09 Apr 2023 07:38:08 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-2g7hz (ro)
Conditions:
  Type          Status
  Initialized   True
  Ready         True
  ContainersReady True
  PodScheduled  True
Volumes:
  kube-api-access-2g7hz:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age    From          Message
  ----     -
  Normal   Scheduled   6m36s  default-scheduler  Successfully assigned default/nginx to aks-agentpool-35627940-vmss000000
  Normal   Pulling     6m36s  kubelet         Pulling image "nginx"
  Normal   Pulled      6m33s  kubelet         Successfully pulled image "nginx" in 2.65517193s
  Normal   Created     6m33s  kubelet         Created container nginx
  Normal   Started     6m33s  kubelet         Started container nginx
```

11. Delete your pod using kubectl delete pod nginx.

```
PS /home/stanislaw> kubectl delete pod nginx
pod "nginx" deleted
PS /home/stanislaw>
```

12. Let's find the image used on one of the coredns pods under the kube-system namespace.

13. Once again list all pods under all namespaces.

```
PS /home/stanislaw> kubectl get pods --all-namespaces
NAMESPACE      NAME                                                    READY   STATUS    RESTARTS   AGE
kube-system    ama-logs-7bnjj                                         2/2     Running   0           20m
kube-system    ama-logs-r77cl                                         2/2     Running   0           20m
kube-system    ama-logs-rs-79b79975d5-5dgmr                         1/1     Running   0           21m
kube-system    azure-ip-masq-agent-m7pjm                             1/1     Running   0           20m
kube-system    azure-ip-masq-agent-nlzh4                             1/1     Running   0           20m
kube-system    cloud-node-manager-rtrnm                             1/1     Running   0           20m
kube-system    cloud-node-manager-zbh2z                             1/1     Running   0           20m
kube-system    coredns-59b6bf8b4f-dcxbv                             1/1     Running   0           21m
kube-system    coredns-59b6bf8b4f-vrnp2                             1/1     Running   0           19m
kube-system    coredns-autoscaler-64b6477b8b-njsnl                  1/1     Running   0           21m
kube-system    csi-azuredisk-node-dnf6r                             3/3     Running   0           20m
kube-system    csi-azuredisk-node-qz4ks                             3/3     Running   0           20m
kube-system    csi-azurefile-node-br9w4                             3/3     Running   0           20m
kube-system    csi-azurefile-node-nwgdz                             3/3     Running   0           20m
kube-system    konnectivity-agent-7cf8bd6556-b7828                 1/1     Running   0           21m
kube-system    konnectivity-agent-7cf8bd6556-t6qqh                 1/1     Running   0           21m
kube-system    kube-proxy-n8q64                                     1/1     Running   0           20m
kube-system    kube-proxy-v52qk                                     1/1     Running   0           20m
kube-system    metrics-server-8655f897d8-glgrl                     2/2     Running   0           7m28s
kube-system    metrics-server-8655f897d8-swgjh                     2/2     Running   0           7m28s
PS /home/stanislaw>
```

14. Note one of the coredns pods. Now run kubectl describe pod <coredns-name> -n kube-system. Replace the

<coredns-name> place holder with noted name.

```
PS /home/stanislaw> kubectl describe pod kube-proxy-v52qk -n kube-system
Name:          kube-proxy-v52qk
Namespace:     kube-system
Priority:       2000001000
Priority Class Name: system-node-critical
Service Account: kube-proxy
Node:          aks-agentpool-35627940-vmss000000/10.224.0.6
Start Time:    Sun, 09 Apr 2023 07:26:25 +0000
Labels:        component=kube-proxy
               controller-revision-hash=6985689589
               pod-template-generation=1
Annotations:   tier=node
Status:        Running
IP:            10.224.0.6
IPs:           IP: 10.224.0.6
Controlled By: DaemonSet/kube-proxy
Init Containers:
kube-proxy-bootstrap:
  Container ID:  containerd://4c89a79a19cb1222ff8d1291c66df263439fb00a10f35b25c796792c2a5e2635
  Image:         mcr.microsoft.com/oss/kubernetes/kube-proxy:v1.24.10.3
  Image ID:      sha256:a03c20def8004c7cca3003be78610ab1b615fd26b91163459b0ef49d08e3b9a
  Port:          <none>
  Host Port:     <none>
  Command:
  /bin/sh
  -c
  SYSCTL=/proc/sys/net/netfilter/nf_conntrack_max
  echo "Current net.netfilter.nf_conntrack_max: $(cat $SYSCTL)"
  DESIRED=$(awk -F= '/net.netfilter.nf_conntrack_max/ {print $2}' /etc/sysctl.d/999-sysctl-aks.conf)
  if [ -z "$DESIRED" ]; then
    DESIRED=$((32768*(nproc)))
    if [ $DESIRED -lt 131072 ]; then
      DESIRED=131072
    fi
  fi
  echo "AKS custom config for net.netfilter.nf_conntrack_max not set."
  echo "Setting nf_conntrack_max to $DESIRED (32768 * $(nproc) cores, minimum 131072)."
  echo $DESIRED > $SYSCTL
else
  echo "AKS custom config for net.netfilter.nf_conntrack_max set to $DESIRED."
  echo "Setting nf_conntrack_max to $DESIRED."
  echo $DESIRED > $SYSCTL
fi
State:          Terminated
```

15. Inspect the output and locate the image information.

16. Now let us check the logs of the metrics-server pod. Run the same command as in step 7 but don't forget to add the namespace in which this pod is created.

```
PS /home/stanislaw> kubectl logs kube-proxy-v52qk -n kube-system -c kube-proxy
I0409 07:26:39.772028 1 flags.go:64] FLAG: --add-dir-header="false"
I0409 07:26:39.772082 1 flags.go:64] FLAG: --alsologtostderr="false"
I0409 07:26:39.772085 1 flags.go:64] FLAG: --bind-address="0.0.0.0"
I0409 07:26:39.772090 1 flags.go:64] FLAG: --bind-address-hard-fail="false"
I0409 07:26:39.772100 1 flags.go:64] FLAG: --boot-id-file="/proc/sys/kernel/random/boot_id"
I0409 07:26:39.772104 1 flags.go:64] FLAG: --cleanup="false"
I0409 07:26:39.772106 1 flags.go:64] FLAG: --cluster-cidr="10.244.0.0/16"
I0409 07:26:39.772110 1 flags.go:64] FLAG: --config=""
I0409 07:26:39.772113 1 flags.go:64] FLAG: --config-sync-period="15m0s"
I0409 07:26:39.772117 1 flags.go:64] FLAG: --conntrack-max-per-core="0"
I0409 07:26:39.772121 1 flags.go:64] FLAG: --conntrack-min="131072"
I0409 07:26:39.772124 1 flags.go:64] FLAG: --conntrack-tcp-timeout-close-wait="1h0m0s"
I0409 07:26:39.772127 1 flags.go:64] FLAG: --conntrack-tcp-timeout-established="24h0m0s"
I0409 07:26:39.772130 1 flags.go:64] FLAG: --detect-local-mode="ClusterCIDR"
I0409 07:26:39.772134 1 flags.go:64] FLAG: --feature-gates=""
I0409 07:26:39.772138 1 flags.go:64] FLAG: --healthz-bind-address="0.0.0.0:10256"
I0409 07:26:39.772141 1 flags.go:64] FLAG: --healthz-port="10256"
I0409 07:26:39.772144 1 flags.go:64] FLAG: --help="false"
I0409 07:26:39.772147 1 flags.go:64] FLAG: --hostname-override=""
I0409 07:26:39.772150 1 flags.go:64] FLAG: --iptables-masquerade-bit="14"
I0409 07:26:39.772153 1 flags.go:64] FLAG: --iptables-min-sync-period="1s"
I0409 07:26:39.772155 1 flags.go:64] FLAG: --iptables-sync-period="30s"
I0409 07:26:39.772158 1 flags.go:64] FLAG: --ipvs-exclude-cidrs="[]"
I0409 07:26:39.772165 1 flags.go:64] FLAG: --ipvs-min-sync-period="0s"
I0409 07:26:39.772169 1 flags.go:64] FLAG: --ipvs-scheduler=""
I0409 07:26:39.772172 1 flags.go:64] FLAG: --ipvs-strict-arp="false"
I0409 07:26:39.772175 1 flags.go:64] FLAG: --ipvs-sync-period="30s"
I0409 07:26:39.772177 1 flags.go:64] FLAG: --ipvs-tcp-timeout="0s"
I0409 07:26:39.772180 1 flags.go:64] FLAG: --ipvs-tcpfin-timeout="0s"
I0409 07:26:39.772183 1 flags.go:64] FLAG: --ipvs-udp-timeout="0s"
I0409 07:26:39.772185 1 flags.go:64] FLAG: --kube-api-burst="10"
I0409 07:26:39.772188 1 flags.go:64] FLAG: --kube-api-content-type="application/vnd.kubernetes.protobuf"
I0409 07:26:39.772192 1 flags.go:64] FLAG: --kube-api-qps="5"
I0409 07:26:39.772196 1 flags.go:64] FLAG: --kubeconfig="/var/lib/kubelet/kubeconfig"
I0409 07:26:39.772199 1 flags.go:64] FLAG: --log-backtrace-at=":0"
I0409 07:26:39.772204 1 flags.go:64] FLAG: --log-dir=""
I0409 07:26:39.772208 1 flags.go:64] FLAG: --log-file=""
I0409 07:26:39.772211 1 flags.go:64] FLAG: --log-file-max-size="1800"
I0409 07:26:39.772213 1 flags.go:64] FLAG: --log-flush-frequency="5s"
I0409 07:26:39.772216 1 flags.go:64] FLAG: --logtostderr="true"
I0409 07:26:39.772219 1 flags.go:64] FLAG: --machine-id-file="/etc/machine-id,/var/lib/dbus/machine-id"
I0409 07:26:39.772222 1 flags.go:64] FLAG: --masquerade-all="false"
I0409 07:26:39.772225 1 flags.go:64] FLAG: --master=""
I0409 07:26:39.772228 1 flags.go:64] FLAG: --metrics-bind-address="0.0.0.0:10249"
I0409 07:26:39.772230 1 flags.go:64] FLAG: --metrics-port="10249"
I0409 07:26:39.772233 1 flags.go:64] FLAG: --nodeport-addresses="[]"
```

Practice2: Working with pod manifest files

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so.

Copy/paste will not help you to learn Kubernetes!

1. Now it is time to deploy pod using manifest file (declarative approach). Copy the following code block on your

local computer in a file called redis.yaml:

```
apiVersion: v1
```

```
kind: pod
```

```
metadata:
```

```
name: static-web
```

```
labels:
```

```
role: myrole
```

```
spec:
```

```
containers:
```

```
- name: redis
```

```
image: redis123
```



- This is the redis.yaml, which I will upload in

GitHub

2. Try to deploy the pod defined in redis.yaml. Run `kubectl create -f redis.yaml`.
3. You will receive errors on your screen. Your next task will be to correct the syntax of the code you just copied. You can use the online Kubernetes documentation or you can search the internet in general.

```
PS C:\Users\stani> kubectl create -f redis.yaml
Unable to connect to the server: dial tcp 127.0.0.1:56386: connectex: No connection could be made because the target machine actively refused it.
PS C:\Users\stani> |
```

4. When you solve all the syntax errors your pod should be deployed but is it running? What is the status of your pod?

```
pod/static-web created
```

5. Check the events associated with this pod. Run the `kubectl describe pod static-web` command. What are the events showing? Why your pod is not running?

| NAME | READY | STATUS | RESTARTS | AGE |
|------------|-------|------------------|----------|-----|
| static-web | 0/1 | ImagePullBackOff | 0 | 27s |

```
Events:
  Type     Reason      Age          From          Message
  ----     -
  Normal   Scheduled   93s         default-scheduler   Successfully assigned default/static-web to docker-desktop
  Warning   Failed      45s (x3 over 91s) kubelet        Failed to pull image "redis123": rpc error: code = Unknown desc = Error response from daemon: pull access denied for redis123, repository does not exist or may require 'docker login': denied: requested access to the resource is denied
  Warning   Failed      45s (x3 over 91s) kubelet        Error: ErrImagePull
  Normal   BackOff     15s (x4 over 90s) kubelet        Back-off pulling image "redis123"
  Warning   Failed      15s (x4 over 90s) kubelet        Error: ImagePullBackOff
  Normal   Pulling     2s (x4 over 93s) kubelet        Pulling image "redis123"
```

6. Find the correct image (check the Docker hub page) and correct it in the manifest.

7. Locate the image information and put the correct image name. Redeploy the pod (first run `kubectl delete pod static-web` to delete the pod, then run `kubectl create` once again).
8. Check the status of your pod. It should be running now.

| NAME | READY | STATUS | RESTARTS | AGE |
|------------|-------|---------|----------|-----|
| static-web | 1/1 | Running | 0 | 7s |

9. Now you can delete the pod. Try to delete it using the `kubectl delete -f redis.yaml`.

```
pod "static-web" deleted
```

10. Your next task is to create and test nginx pod definition. Your definition should use the nginx official image, should use label named app with value frontend and should publish port 80. Make sure you complete this task because we will use this template in our next Labs. Your nginx pod should be running without any issues.



- the yaml file.

| Events: | | | | |
|---------|-----------|-----|-------------------|-----------------------------------------------------------|
| Type | Reason | Age | From | Message |
| Normal | Scheduled | 77s | default-scheduler | Successfully assigned default/nginx-pod to docker-desktop |
| Normal | Pulling | 77s | kubelet | Pulling image "nginx" |
| Normal | Pulled | 63s | kubelet | Successfully pulled image "nginx" in 13.57627594s |
| Normal | Created | 63s | kubelet | Created container nginx-container |
| Normal | Started | 63s | kubelet | Started container nginx-container |

11. Final task of this practice will be to define pod definition with following details:

- The yaml file for nginx pod.
- Image=memcached
- Port= 11211
- Label app=web
- CPU request=0.35 cores
- RAM request=0.15 GB
- CPU limit=0.5 cores
- Ram limit=0.25 GB
- Restart policy=Never

12. Don't forget to try your pod definition.



memcached-web.yaml

- the yaml file for the pod.

```
pod/memcached-web created
```

| NAME | READY | STATUS | RESTARTS | AGE |
|---------------|-------|---------|----------|------|
| memcached-web | 1/1 | Running | 0 | 15s |
| nginx-pod | 1/1 | Running | 0 | 7m7s |

Practice3: Multi-container pods

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so. Copy/paste will not help you to learn Kubernetes!

1. Once finished you can try to create multi-container pod definition. Your multi-container pod should use redis and nginx containers with port 6379 and 80 published respectively. Label name should be app with value web.



- the yaml file for the 2 containers in 1 pod.

2. Note that in reality there is no sense to put the redis and nginx under the same pod but it can be done for the purpose of learning.

3. Deploy your multi-container pod. It should have running status. What is written under Ready column when you kubectl get the pods? Why your pod displays different values for ready?

| | | | | |
|--------|-----|---------|---|-----|
| webapp | 2/2 | Running | 0 | 66s |
|--------|-----|---------|---|-----|

4. Kubectl describe you new pod, and locate the containers section. How many containers are listed?

- There are 2 containers, running in one pod.

```
Containers:
  redis:
    Container ID:   docker://e2630b2c8b28c2339972f5dfbddd6d77e9cbbcc36e908e59bc1e4238589e58a
    Image:          asterixlegaulois/redis123
    Image ID:       docker-pullable://asterixlegaulois/redis123@sha256:cb2ddf11373c66e8c66ea7cb4
                    ddd82bd69849d4838ae41ee0e71a0bc5c6cc4c4
    Port:          6379/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 04 Apr 2023 13:08:44 +0200
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-8zgfz (ro)
  nginx-container:
    Container ID:   docker://12a72152b265220b56f0bc81f61f78265c0c97a76a53c04a94aea93ee4833f5a
    Image:          nginx
    Image ID:       docker-pullable://nginx@sha256:2ab30d6ac53580a6db8b657abf0f68d75360ff5cc1670
                    a85acb5bd85ba1b19c0
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 04 Apr 2023 13:08:45 +0200
    Ready:         True
```

5. Delete all the pods under the default namespace.

```
pod "memcached-web" deleted
pod "nginx-pod" deleted
pod "webapp" deleted
```

6. Don't delete any of the manifest files you have created so far.

- They all will be in my Github.

Practice4: Probes

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so. Copy/paste will not help you to learn Kubernetes!

1. First we will create and test liveness probe with exec test. Create a file named probes_exec.yaml with following content:



liveness-exec.yaml

- The yaml file for the container

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  labels:
```

```
    test: liveness
```

```
  name: liveness-exec
```

```
spec:
```

```
  containers:
```

```
    - name: liveness
```

```
      image: k8s.gcr.io/busybox
```

```
      args:
```

```
        - /bin/sh
```

```
        - -c
```

- touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600

livenessProbe:

exec:

command:

- cat

- /tmp/healthy

initialDelaySeconds: 5

periodSeconds: 5

2. Examine the containers args commands especially the line that start with touch. This bash pipeline will help us to test the liveness probes.

3. Run `kubectl create -f probes_exec.yaml`.

```
pod/liveness-exec created
```

4. Run `kubectl describe pod liveness-exec` immediately after you deploy the pod. The output should indicate that no liveness probes have failed yet.

```
Events:
  Type    Reason      Age   From              Message
  ----    -
  Normal  Scheduled   15s   default-scheduler Successfully assigned default/liveness-exec to doc
ker-desktop
  Normal  Pulling     14s   kubelet           Pulling image "k8s.gcr.io/busybox"
  Normal  Pulled      12s   kubelet           Successfully pulled image "k8s.gcr.io/busybox" in
2.11243698s
  Normal  Created     12s   kubelet           Created container liveness
  Normal  Started     12s   kubelet           Started container liveness
```

5. After 35 seconds, view the Pod events again. Run `kubectl describe pod liveness-exec`.

6. At the bottom of the output, there should be a messages indicating that the liveness probes have failed, and the containers have been killed and recreated.

```
Events:
  Type      Reason      Age           From          Message
  ----      -
  Normal    Scheduled   68s          default-scheduler   Successfully assigned default/liveness-exec to docker-desktop
  Normal    Pulling     67s          kubelet        Pulling image "k8s.gcr.io/busybox"
  Normal    Pulled      65s          kubelet        Successfully pulled image "k8s.gcr.io/busybox" in 2.11243698s
  Normal    Created     65s          kubelet        Created container liveness
  Normal    Started     65s          kubelet        Started container liveness
  Warning   Unhealthy   23s (x3 over 33s) kubelet        Liveness probe failed: cat: can't open '/tmp/healthy': No such file or directory
  Normal    Killing     23s          kubelet        Container liveness failed liveness probe, will be restarted
```

7. Wait another 30 seconds, and verify that the container has been restarted. Run `kubectl get pod liveness-exec`.

```
PS C:\Users\V&M\Desktop\lab> kubectl get pod liveness-exec
NAME          READY   STATUS    RESTARTS   AGE
liveness-exec 1/1     Running   1 (55s ago) 2m10s
PS C:\Users\V&M\Desktop\lab> _
```

8. The output should show that RESTARTS has been incremented.

9. We will continue with HTTP probe. Create file named `probes_http.yaml` with following content:



`probes_http.yaml`

- The yaml file for this containers.

`apiVersion: v1`

`kind: Pod`

`metadata:`

labels:

test: liveness

name: liveness-http

spec:

containers:

- name: liveness

image: k8s.gcr.io/liveness

args:

- /server

livenessProbe:

httpGet:

path: /healthz

port: 8080

httpHeaders:

- name: Custom-Header

value: Awesome

initialDelaySeconds: 3

periodSeconds: 3

10. Just for your info, /healthz handler has following function implemented:

```
http.HandleFunc("/healthz", func(w http.ResponseWriter, r
*http.Request) {
```

```

duration := time.Now().Sub(started)
if duration.Seconds() > 10 {
w.WriteHeader(500)
w.Write([]byte(fmt.Sprintf("error: %v", duration.Seconds())))
} else {
w.WriteHeader(200)
w.Write([]byte("ok"))
}
})

```

11. For the first 10 seconds that the container is alive, the /healthz handler returns a status of 200.

After that, the handler returns a status of 500.

12. Run `kubectl create -f probes_http.yaml`.

```
pod/liveness-http created
```

13. Immediately run (you only have 10 secs to run this command) `kubectl describe pod liveness-http`.

```

Name:          liveness-http
Namespace:     default
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.4
Start Time:    Tue, 04 Apr 2023 13:37:15 +0200
Labels:        test=liveness
Annotations:    <none>
Status:        Running
IP:            10.1.0.18

```

14. Your pod should be live and running.

15. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted.

Run again `kubectl describe pod liveness-http`.

```
Events:
  Type       Reason      Age           From          Message
  ----       -
  Normal     Scheduled   80s          default-scheduler   Successfully assigned default-scheduler/liveness-http to docker-desktop
  Normal     Pulled      78s          kubelet        Successfully pulled image "k8s.gcr.io/liveness" in 1.216349349s
  Normal     Pulled      60s          kubelet        Successfully pulled image "k8s.gcr.io/liveness" in 1.144994584s
  Normal     Created     42s (x3 over 78s) kubelet        Created container liveness
  Normal     Started     42s (x3 over 78s) kubelet        Started container liveness
  Normal     Pulled      42s          kubelet        Successfully pulled image "k8s.gcr.io/liveness" in 1.182641358s
  Normal     Pulling     26s (x4 over 79s) kubelet        Pulling image "k8s.gcr.io/liveness"
  Warning    Unhealthy   26s (x9 over 68s) kubelet        Liveness probe failed: HTTP probe failed with statuscode: 500
  Normal     Killing     26s (x3 over 62s) kubelet        Container liveness failed liveness probe, will be restarted
```

16. You should see the same output as in step 7. Kubelet will reboot the container.

17. We continue with TCP probes. Create file named `probes_tcp.yaml` with following content:



`probes_tcp.yaml`

-The yaml file for this containers

`apiVersion: v1`

`kind: Pod`

`metadata:`

`name: liveness-tcp`

`labels:`

app: goproxy

spec:

containers:

- name: goproxy

image: k8s.gcr.io/goproxy:0.1

ports:

- containerPort: 8080

livenessProbe:

tcpSocket:

port: 9999 #8080 is valid port

initialDelaySeconds: 15

periodSeconds: 20

18. Run `kubectl create -f probes_tcp.yaml`.

```
pod/liveness-tcp created
```

19. Immediately run (you only have 10 secs to run this command) `kubectl describe pod liveness-tcp`.

```
Name:          liveness-tcp
Namespace:     default
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.4
Start Time:    Tue, 04 Apr 2023 13:58:20 +0200
Labels:        app=goproxy
Annotations:    <none>
Status:        Running
IP:            10.1.0.19
```

20. Your pod should be live and running.

21. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again `kubectl describe pod liveness-tcp`.

| Events: | | | | | |
|---------|-----------|-------------------|-------------------|--------------------------------------------------------------------|--|
| Type | Reason | Age | From | Message | |
| Normal | Scheduled | 74s | default-scheduler | Successfully assigned default/liveness-tcp to docker-desktop | |
| Normal | Pulling | 73s | kubelet | Pulling image "k8s.gcr.io/goproxy:0.1" | |
| Normal | Pulled | 71s | kubelet | Successfully pulled image "k8s.gcr.io/goproxy:0.1" in 2.233846916s | |
| Normal | Created | 13s (x2 over 70s) | kubelet | Created container goproxy | |
| Normal | Started | 13s (x2 over 70s) | kubelet | Started container goproxy | |
| Warning | Unhealthy | 13s (x3 over 53s) | kubelet | Liveness probe failed: dia | |

22. You should see the same output as in step 7 and 16.

Kubelet will reboot the container.

23. Our last job will be to define one readiness probe using HTTP test.

24. Create file named `readiness_http.yaml` with following content:



`readiness.yaml`

-The last yaml file for this container

`apiVersion: v1`

`kind: Pod`

`metadata:`

`name: readiness-http`

`labels:`

`app: test`

`spec:`

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

readinessProbe:

initialDelaySeconds: 1

periodSeconds: 2

timeoutSeconds: 1

successThreshold: 1

failureThreshold: 1

httpGet:

host:

scheme: HTTP

path: /

httpHeaders:

- name: Host

value: myapplication1.com

port: 80

25. Run `kubectl create -f readiness_http.yaml`.

```
pod/readiness-http created
```

26. Run `kubectl get pods -A` to see the status of your pod.

| NAMESPACE | NAME | READY | STATUS | RESTARTS | AGE |
|-------------|----------------------------------------|-------|------------------|----------------|-------|
| default | liveness-exec | 0/1 | CrashLoopBackOff | 13 (56s ago) | 37m |
| default | liveness-http | 0/1 | CrashLoopBackOff | 13 (3m46s ago) | 28m |
| default | liveness-tcp | 0/1 | CrashLoopBackOff | 5 (70s ago) | 7m11s |
| default | readiness-http | 1/1 | Running | 0 | 7s |
| kube-system | coredns-565d847f94-f45t7 | 1/1 | Running | 0 | 120m |
| kube-system | coredns-565d847f94-sjcxj | 1/1 | Running | 0 | 120m |
| kube-system | etcd-docker-desktop | 1/1 | Running | 1 | 120m |
| kube-system | kube-apiserver-docker-desktop | 1/1 | Running | 1 | 120m |
| kube-system | kube-controller-manager-docker-desktop | 1/1 | Running | 1 | 120m |
| kube-system | kube-proxy-6h7fw | 1/1 | Running | 0 | 120m |
| kube-system | kube-scheduler-docker-desktop | 1/1 | Running | 1 | 120m |
| kube-system | storage-provisioner | 1/1 | Running | 0 | 119m |
| kube-system | vpnkit-controller | 1/1 | Running | 9 (7m8s ago) | 119m |

27. Pods and their status and ready states will be displayed; our pod should be in running state.

28. Run `kubectl describe pod readiness-http`. Examine the events for this pod. Everything should be OK.

29. Now delete the pod and edit the `readiness_http.yaml` so that the port parameter has 81 value.

30. Run again `kubectl create -f readiness_http.yaml`.

31. Run `kubectl get pods -A` to see the status of your pod. You should see that the pod is running but it is not in ready state.

| NAMESPACE | NAME | READY | STATUS | RESTARTS | AGE |
|-------------|----------------------------------------|-------|------------------|----------------|------|
| default | liveness-exec | 0/1 | CrashLoopBackOff | 13 (4m46s ago) | 41m |
| default | liveness-http | 0/1 | CrashLoopBackOff | 15 (111s ago) | 32m |
| default | liveness-tcp | 0/1 | CrashLoopBackOff | 7 (80s ago) | 11m |
| default | readiness-http | 0/1 | Running | 0 | 10s |
| kube-system | coredns-565d847f94-f45t7 | 1/1 | Running | 0 | 123m |
| kube-system | coredns-565d847f94-sjcxj | 1/1 | Running | 0 | 123m |
| kube-system | etcd-docker-desktop | 1/1 | Running | 1 | 123m |
| kube-system | kube-apiserver-docker-desktop | 1/1 | Running | 1 | 124m |
| kube-system | kube-controller-manager-docker-desktop | 1/1 | Running | 1 | 123m |
| kube-system | kube-proxy-6h7fw | 1/1 | Running | 0 | 123m |
| kube-system | kube-scheduler-docker-desktop | 1/1 | Running | 1 | 124m |
| kube-system | storage-provisioner | 1/1 | Running | 0 | 123m |
| kube-system | vpnkit-controller | 1/1 | Running | 10 (2m55s ago) | 123m |

32. Describe the pod. Run `kubectl describe pod readiness-http`.

33. From the events we can see that readiness probe failed due to the connection being refused therefore pod will not receive any traffic.

34. Delete all pods under the default namespace.

```
pod "liveness-exec" deleted  
pod "liveness-http" deleted  
pod "liveness-tcp" deleted  
pod "readiness-http" deleted
```

35. Don't delete any manifest files created so far